Project 3

You are hired as a DevOps Engineer for Analytics Pvt Ltd. This company is a product-based organization which uses Docker for their containerization needs within the company. The final product received a lot of traction in the first few weeks of launch. Now with the increasing demand, the organization needs to have a platform for automating deployment, scaling and operations of application containers across clusters of hosts. As a DevOps Engineer, you need to implement a DevOps lifecycle such that all the requirements are implemented without any change in the Docker containers in the testing environment.

Up until now, this organization used to follow a monolithic architecture with just 2 developers. The product is present on: https://github.com/hshar/website.git Following are the specifications of the lifecycle:

1. Git workflow should be implemented. Since the company follows a monolithic architecture of development, you need to take care of version control. The release should happen only on the 25th of every month.
2. CodeBuild should be triggered once the commits are made in the master branch.
3. The code should be containerized with the help of the Dockerfile. The Dockerfile should be built every time if there is a push to GitHub. Create a custom Docker image using a Dockerfile.
4. As per the requirement in the production server, you need to use the Kubernetes cluster and the containerized code from Docker Hub should be deployed with 2 replicas. Create a NodePort service and configure the same for port 30008.
5. Create a Jenkins Pipeline script to accomplish the above task.
6. For configuration management of the infrastructure, you need to deploy the configuration on the servers to install necessary software and configurations.
7. Using Terraform, accomplish the task of infrastructure creation in the AWS cloud provider.

Architectural Advice:

Software's to be installed on the respective machines using configuration management.

Worker1: Jenkins, Java

Worker2: Docker, Kubernetes

Worker3: Java, Docker, Kubernetes

Worker4: Docker, Kubernetes

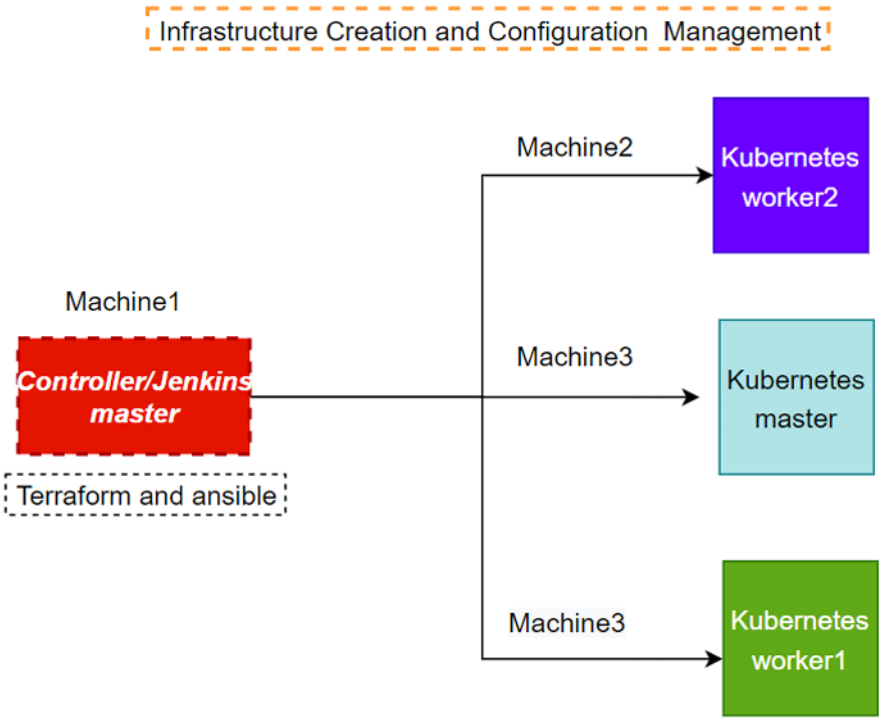

Servers for jenkins and kubernetes configuration

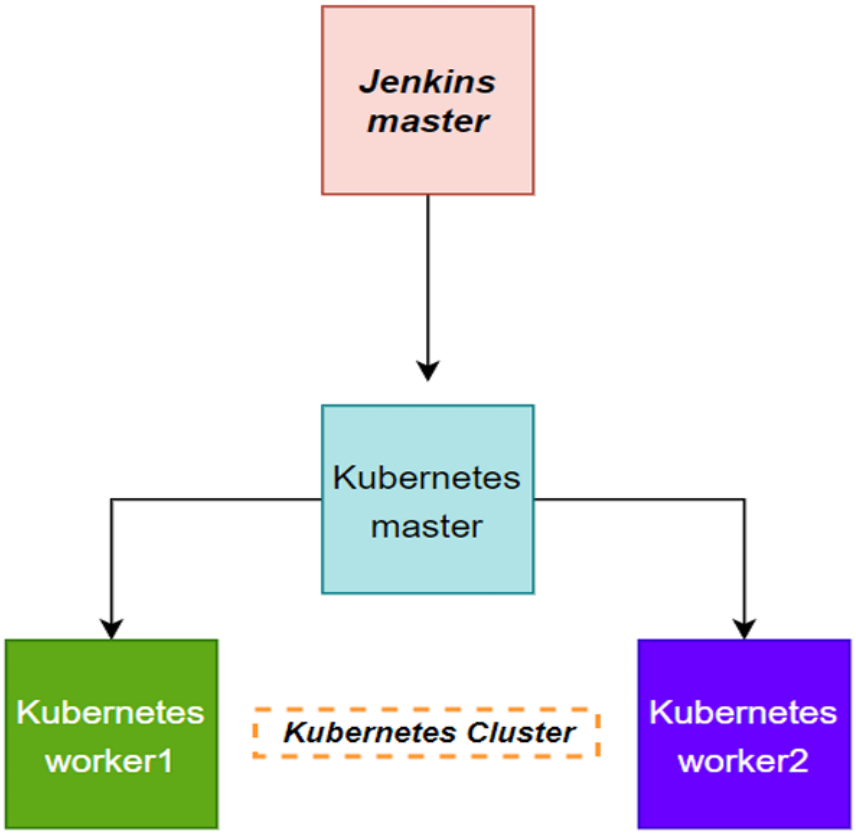

DevOps Project

Architecture

We have total 4 Machine. Machine 1, 2, 3 and 4

- Machine-1 - Terraform, Ansible, Java and Jenkins Installed
- Machine-2 – Docker and Kubernetes
- Machine-3 – Java, Docker and Kubernetes
- Machine-4 - Docker and Kubernetes

Machine 1 (Jenkins Master)→Machine 2(Kubernetes Master)→ Machine 3 and Machine 4

1. Creating Machine-1

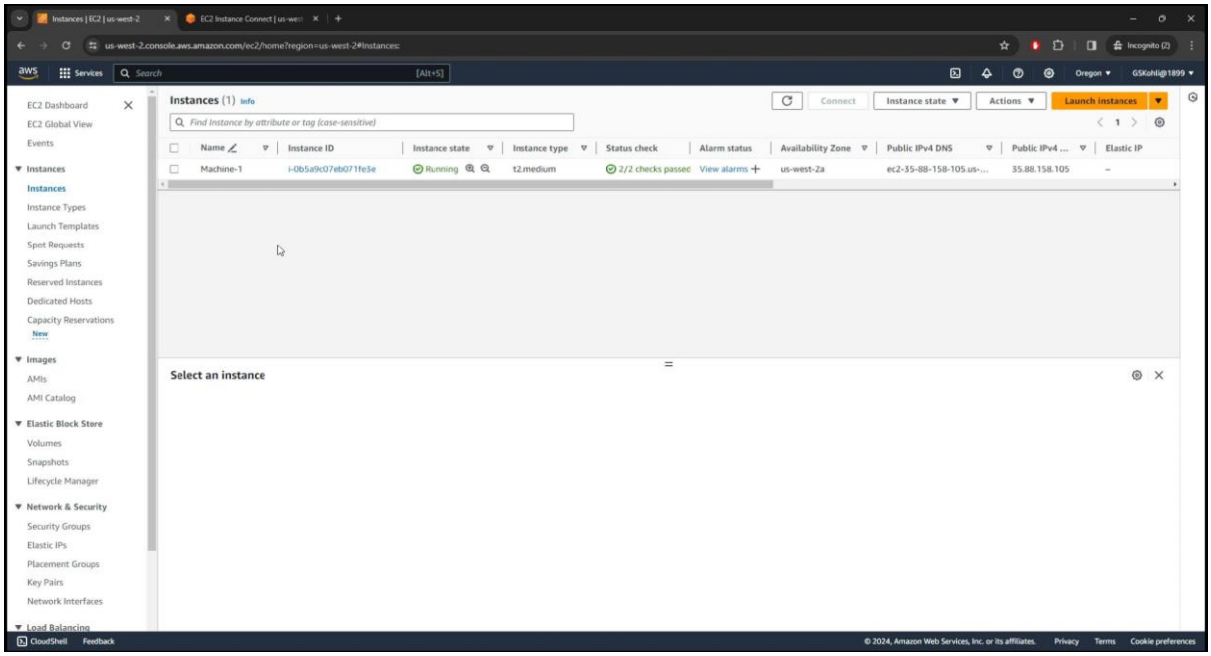

2. Making script to install required packages for Machine 1

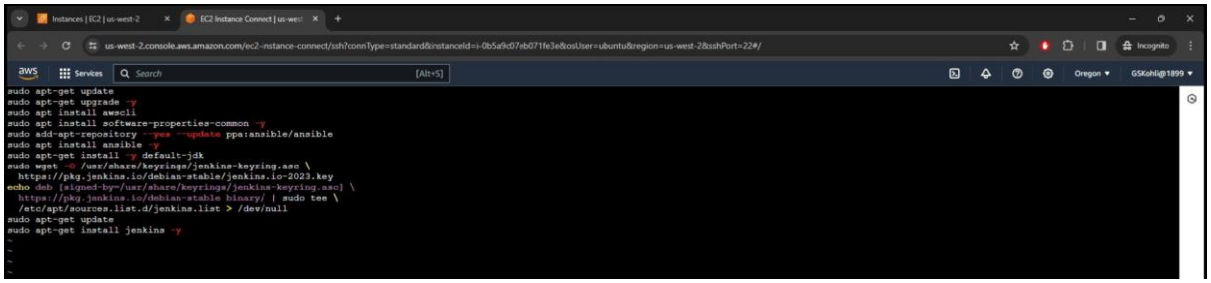

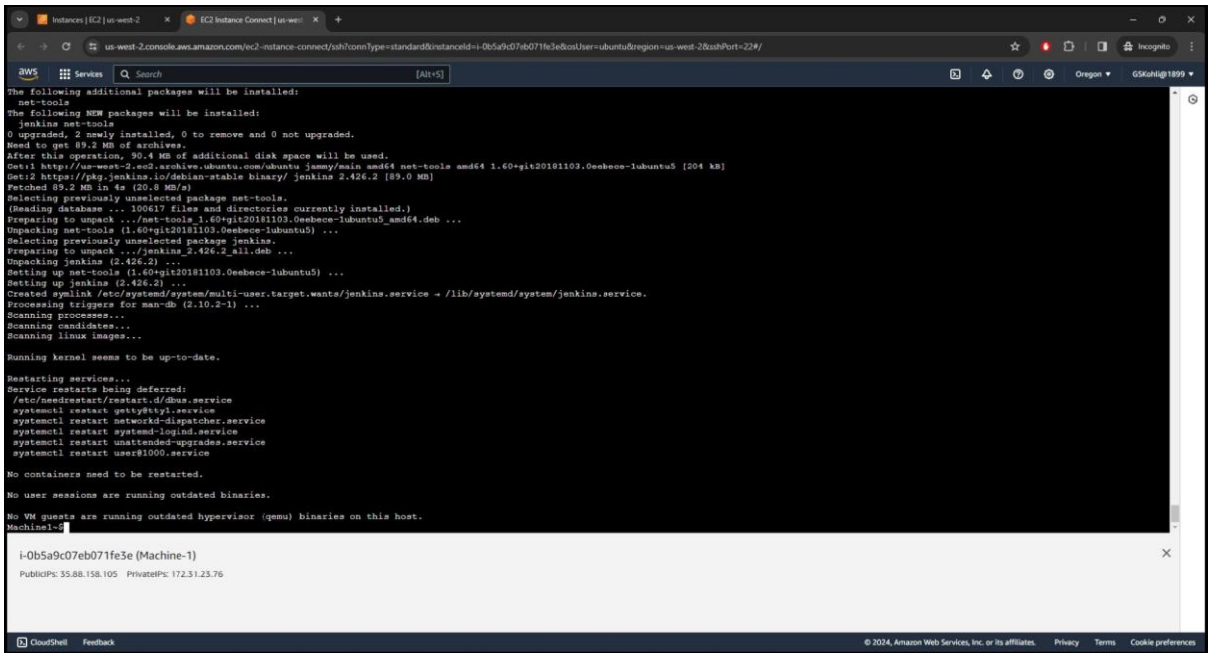

3. Making a terraform script to create all other Machines.

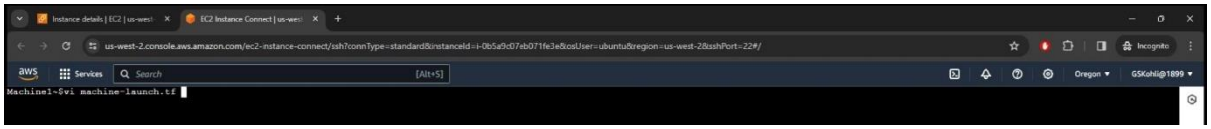

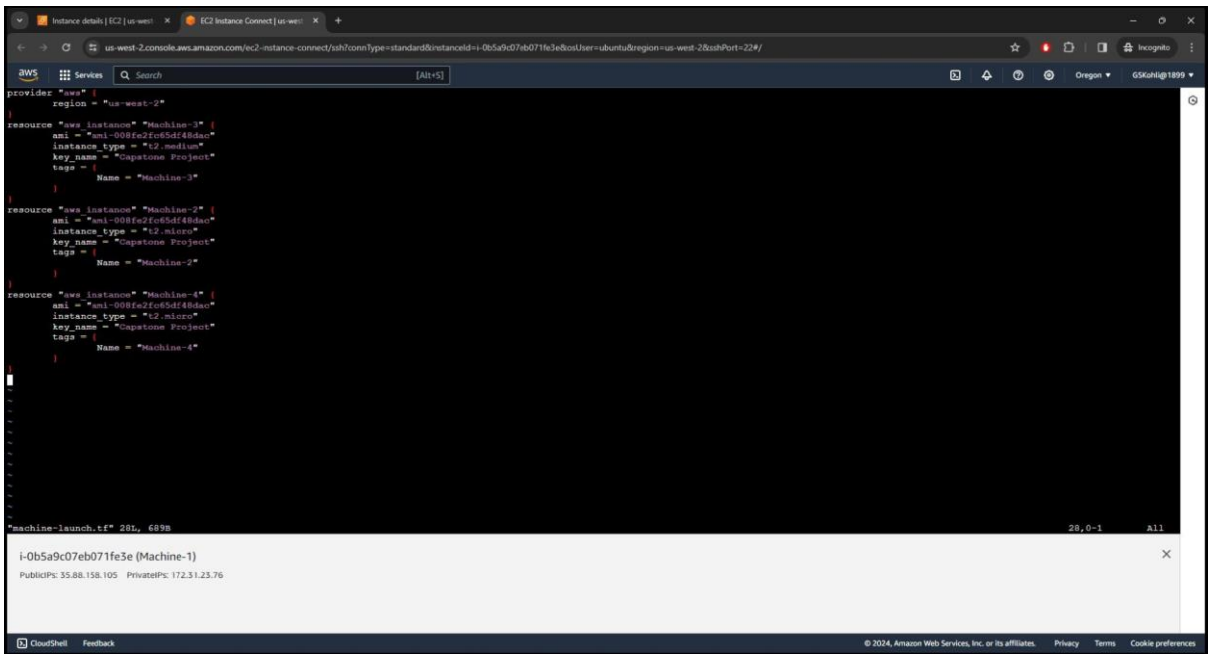

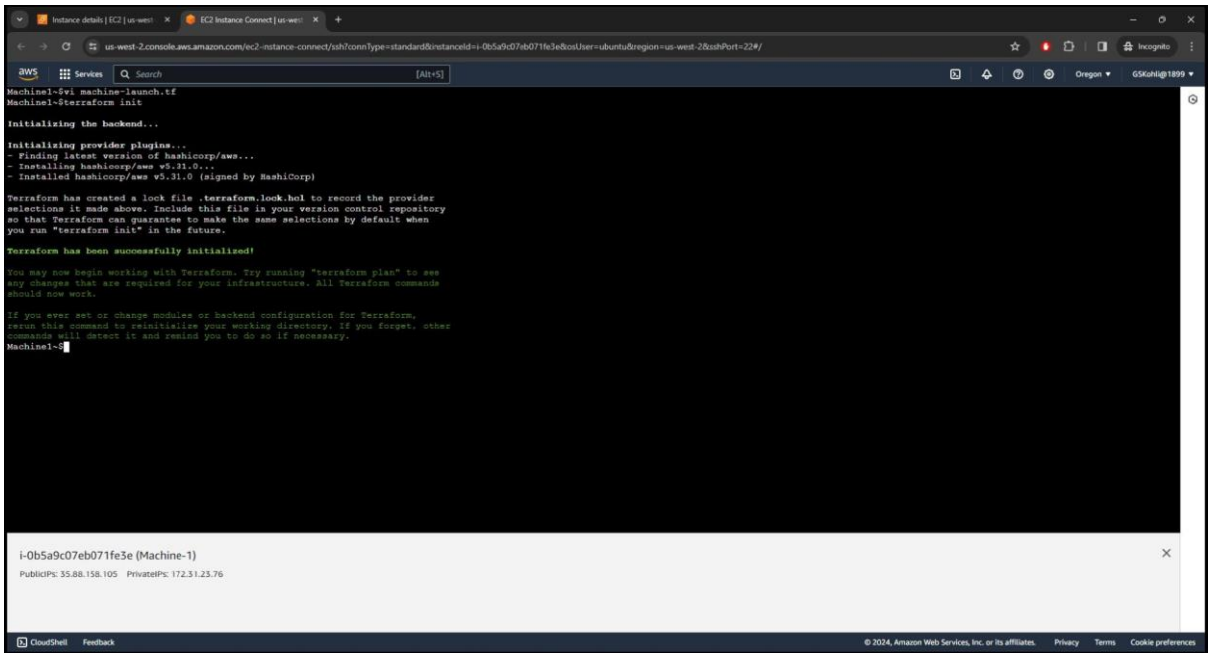
Machine1~$vi machine-launch.tf
Machine1~$terraform init
Initializing the backend...
Initializing provider plugins...
- Finding latest version of hashicorp/aws...
- Installing hashicorp/aws v5.31.0...
- Installed hashicorp/aws v5.31.0 (signed by HashiCorp)
Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.
Terraform has been successfully initialized!
You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.
If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
Machine1~$
i-0b5a9c07eb071fe3e (Machine-1)
PublicIPs: 35.88.158.105 PrivateIPs: 172.31.23.76

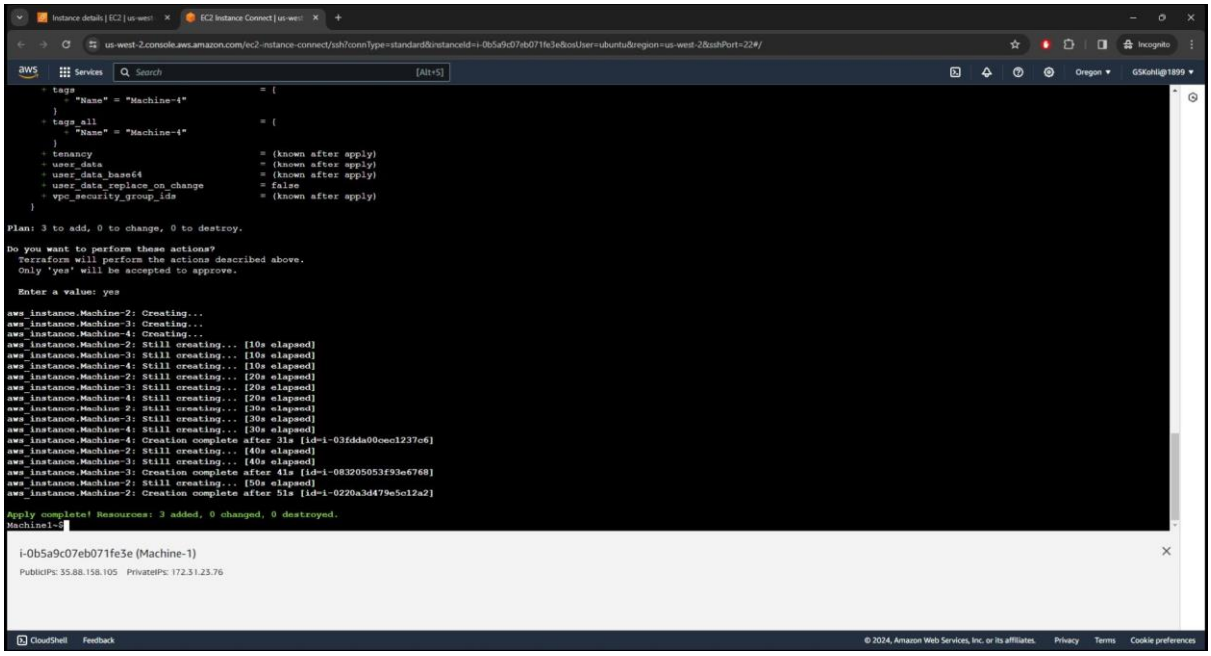
+ tags = {
+ "Name" = "Machine-4"
}
+ tags_all = {
+ "Name" = "Machine-4"
}
+ tenancy = (known after apply)
+ user_data = (known after apply)
+ user_data_base64 = (known after apply)
+ user_data_replace_on_change = false
+ vpc_security_group_ids = (known after apply)
}
Plan: 3 to add, 0 to change, 0 to destroy.
Do you want to perform these actions?
Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.
Enter a value: yes
aws_instance.Machine-2: Creating...
aws_instance.Machine-3: Creating...
aws_instance.Machine-4: Creating...
aws_instance.Machine-2: Still creating... [10s elapsed]
aws_instance.Machine-3: Still creating... [10s elapsed]
aws_instance.Machine-4: Still creating... [10s elapsed]
aws_instance.Machine-2: Still creating... [20s elapsed]
aws_instance.Machine-3: Still creating... [20s elapsed]
aws_instance.Machine-4: Still creating... [20s elapsed]
aws_instance.Machine-2: Still creating... [30s elapsed]
aws_instance.Machine-3: Still creating... [30s elapsed]
aws_instance.Machine-4: Still creating... [30s elapsed]
aws_instance.Machine-4: Creation complete after 31s [id=i-03fdda00cec1237c6]
aws_instance.Machine-2: Still creating... [40s elapsed]
aws_instance.Machine-3: Still creating... [40s elapsed]
aws_instance.Machine-3: Creation complete after 41s [id=i-083205053f93e6768]
aws_instance.Machine-2: Still creating... [50s elapsed]
aws_instance.Machine-2: Creation complete after 51s [id=i-0220a3d479e5c12a2]
Apply complete! Resources: 3 added, 0 changed, 0 destroyed.
Machine1~$
i-0b5a9c07eb071fe3e (Machine-1)
PublicIPs: 35.88.158.105 PrivateIPs: 172.31.23.76

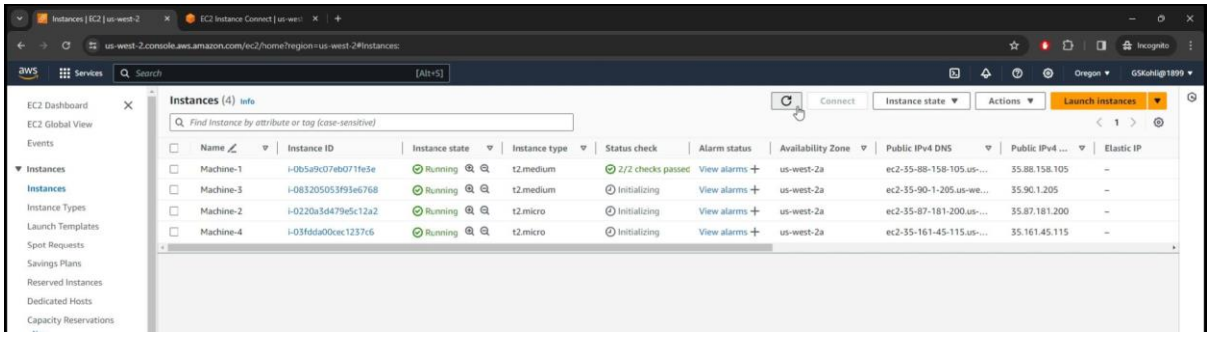
Instances (4)
Name
Instance ID
Instance state
Instance type
Status check
Alarm status
Availability Zone
Public IPv4 DNS
Public IPv4 ...
Elastic IP
Machine-1 i-0b5a9c07eb071fe3e Running t2.medium 2/2 checks passed us-west-2a ec2-35-88-158-105.us-... 35.88.158.105
Machine-3 i-083205053f93e6768 Running t2.medium Initializing us-west-2a ec2-35-90-1-205.us-we... 35.90.1.205
Machine-2 i-0220a3d479e5c12a2 Running t2.micro Initializing us-west-2a ec2-35-87-181-200.us-... 35.87.181.200
Machine-4 i-03fdda00cec1237c6 Running t2.micro Initializing us-west-2a ec2-35-161-45-115.us-... 35.161.45.115

4. Setting up Ansible

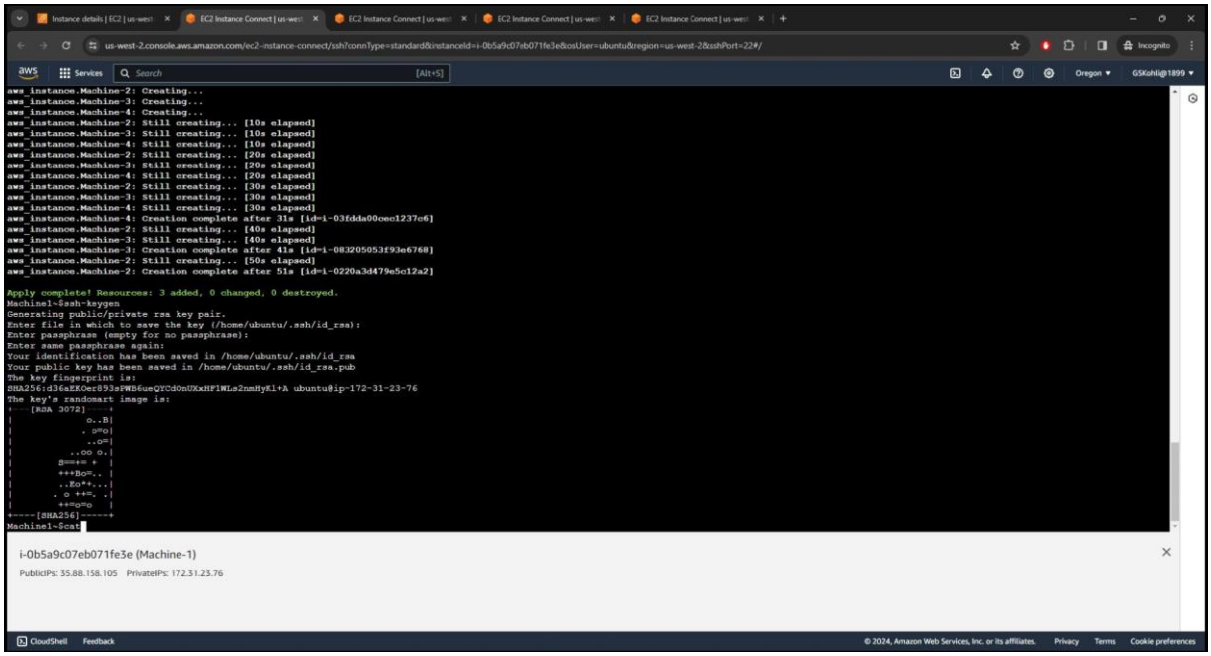

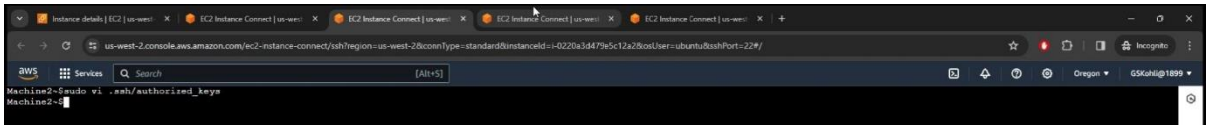

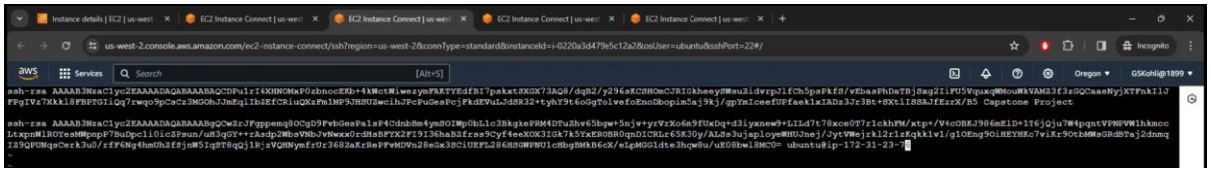

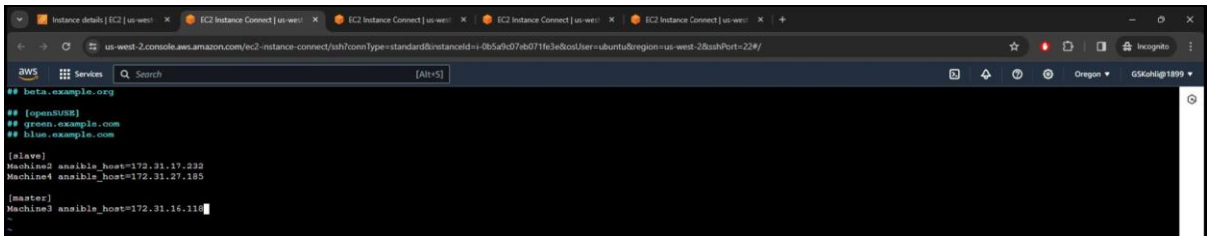

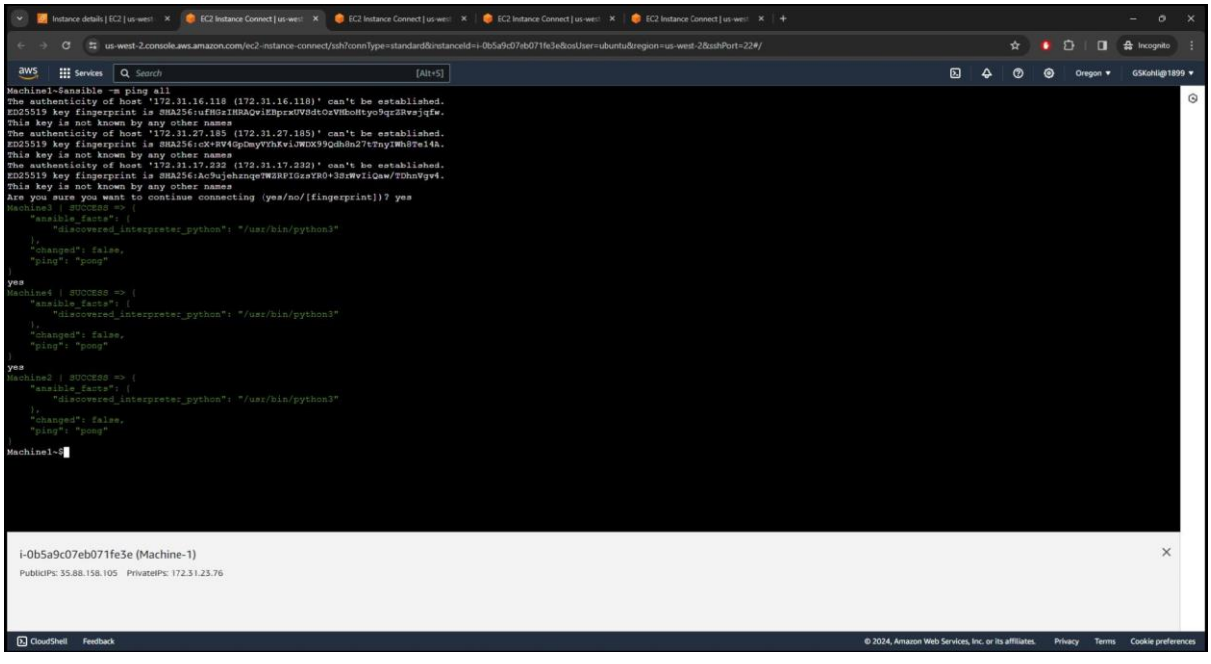

5. Making 2 shell script and 1 yaml script to install required packages in machine 2, 3 and 4

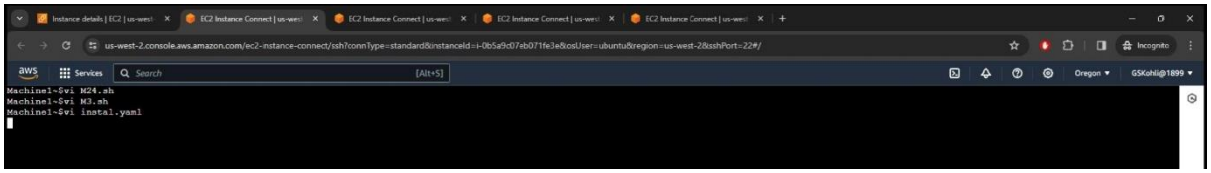

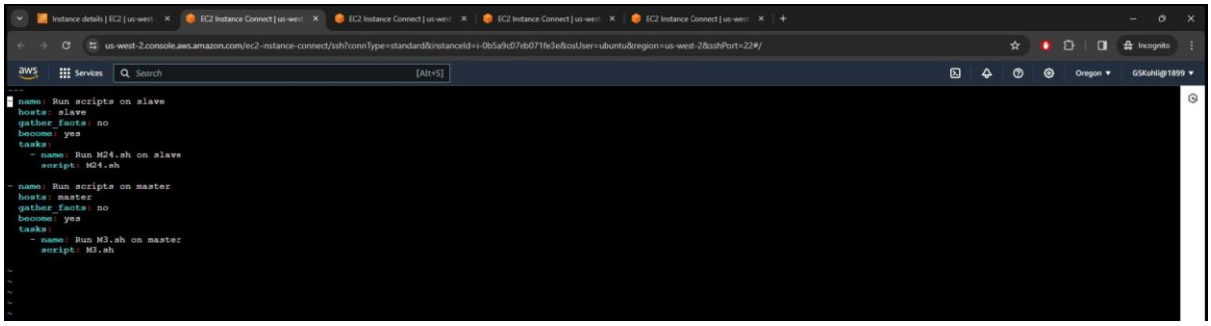

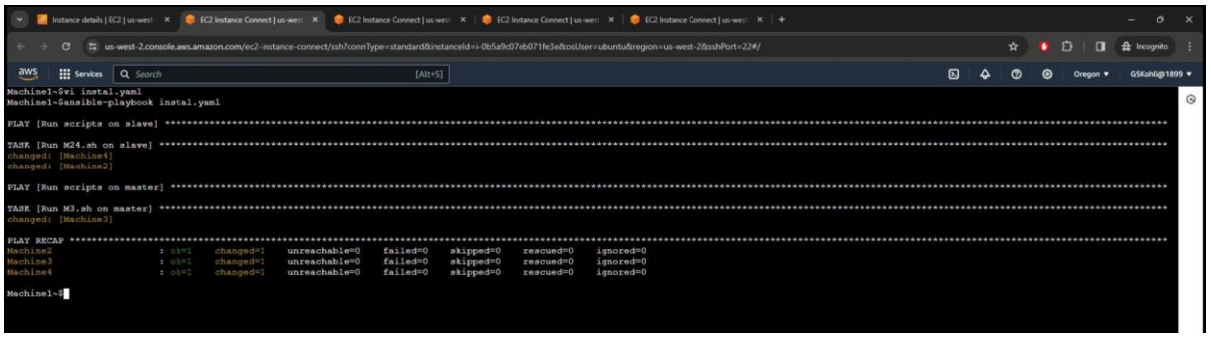

6. Setting Kubernetes Cluster

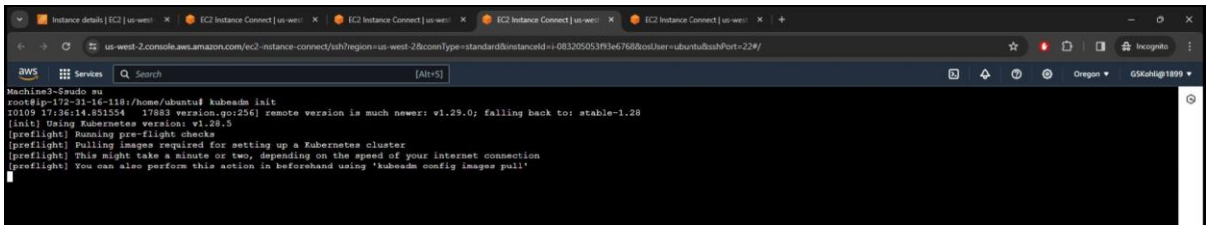

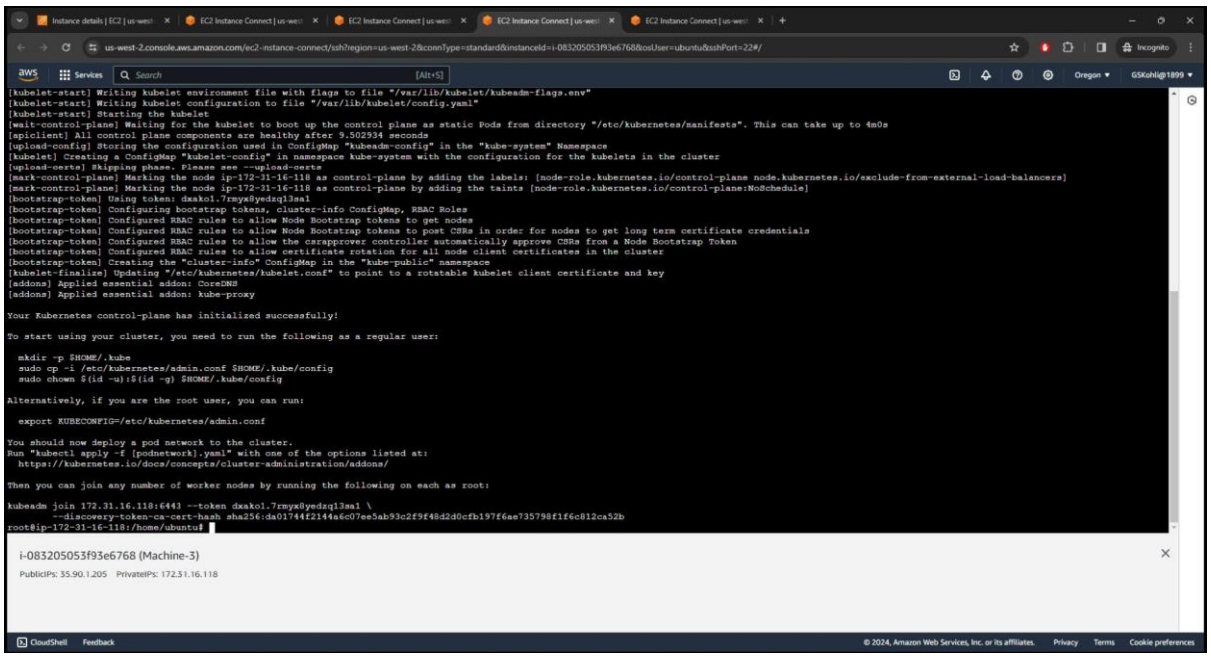

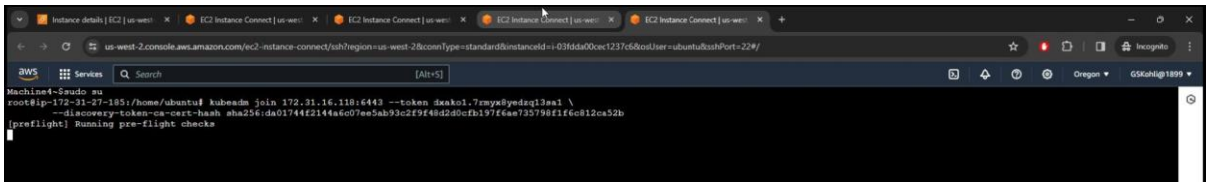

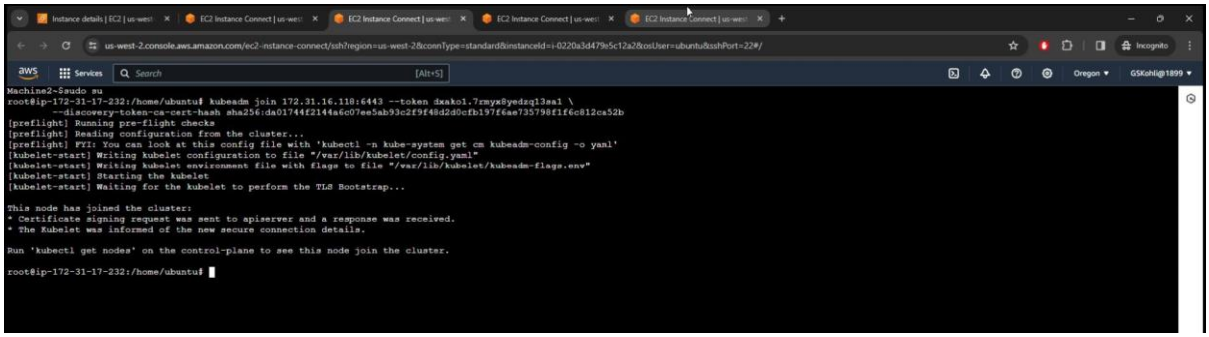

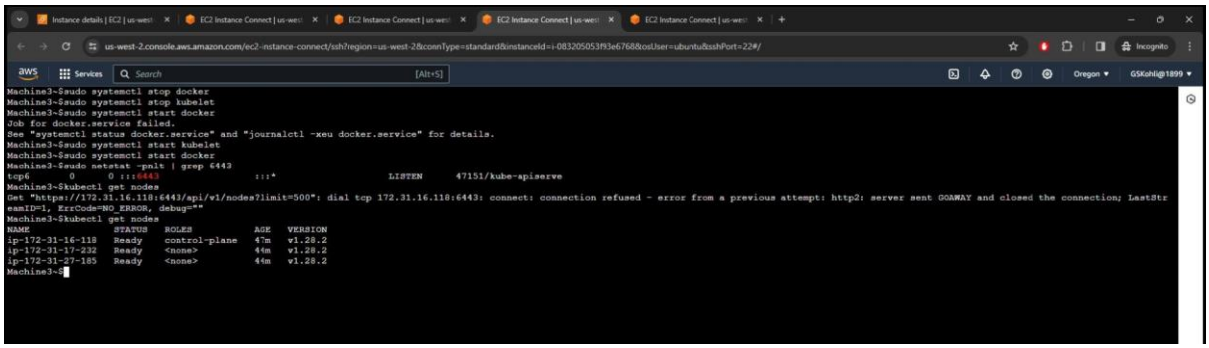

7. Now Setting Jenkins

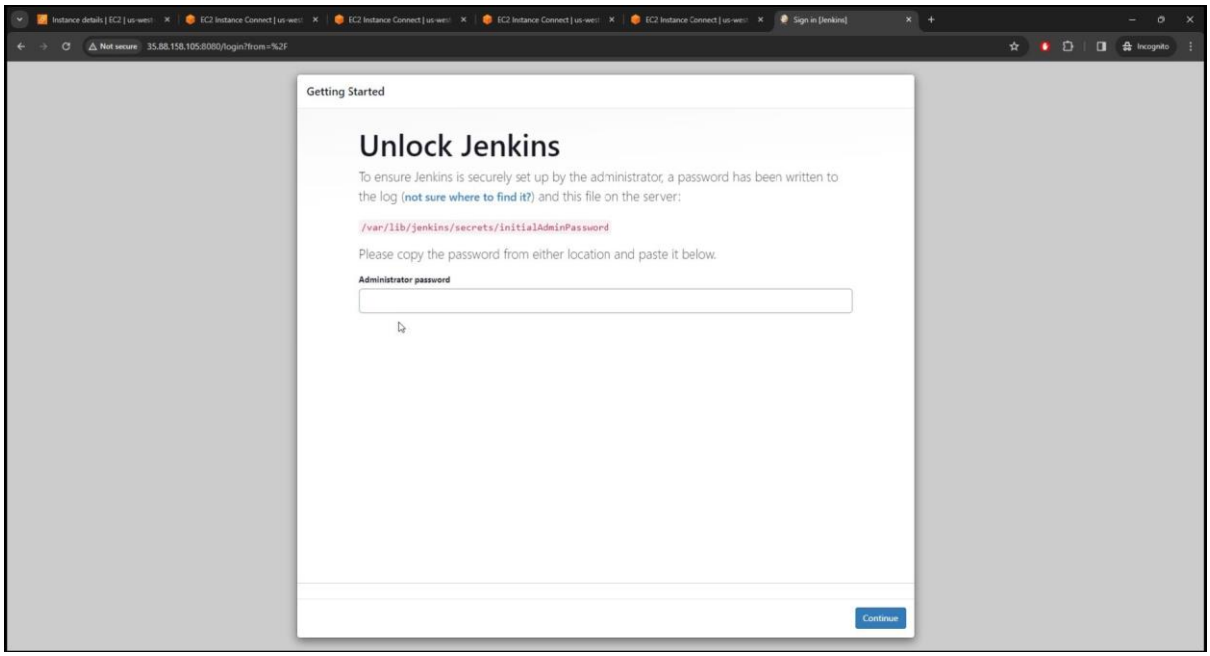

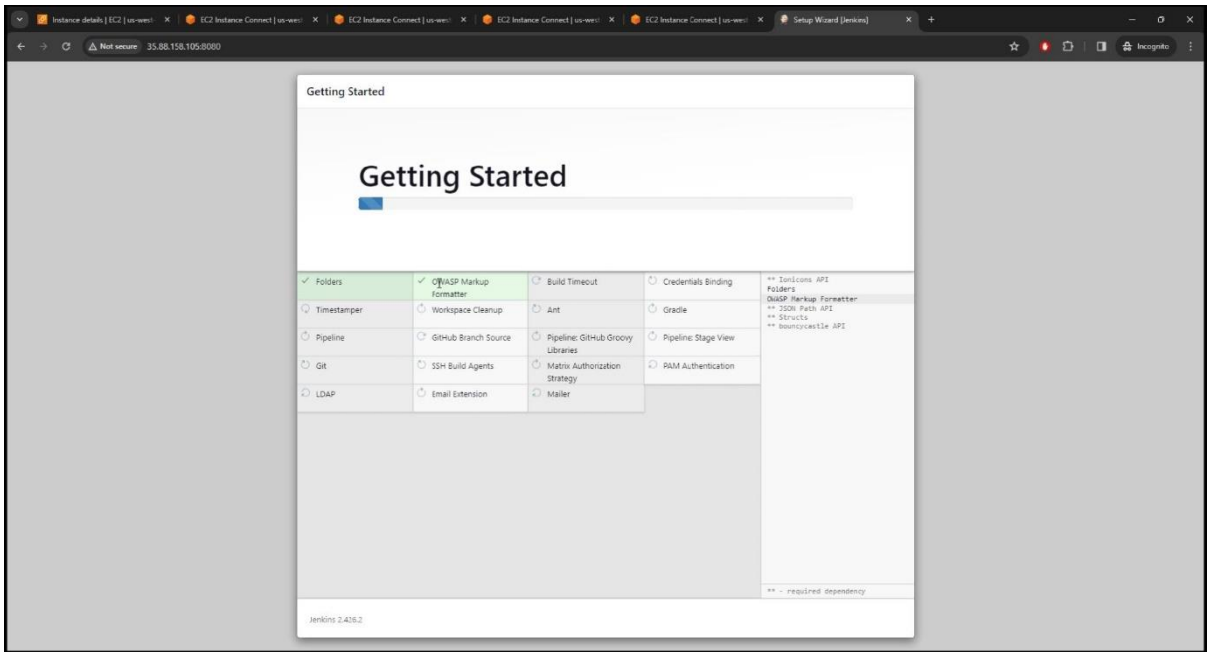

8. Adding Machine 3 as a node.

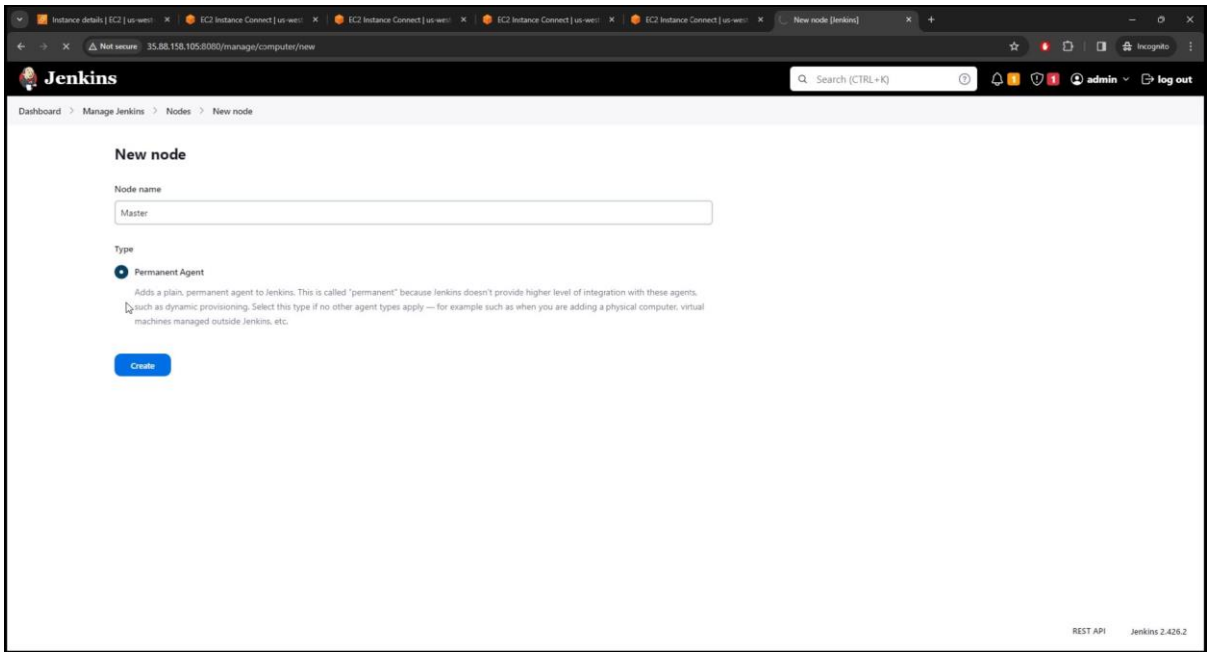

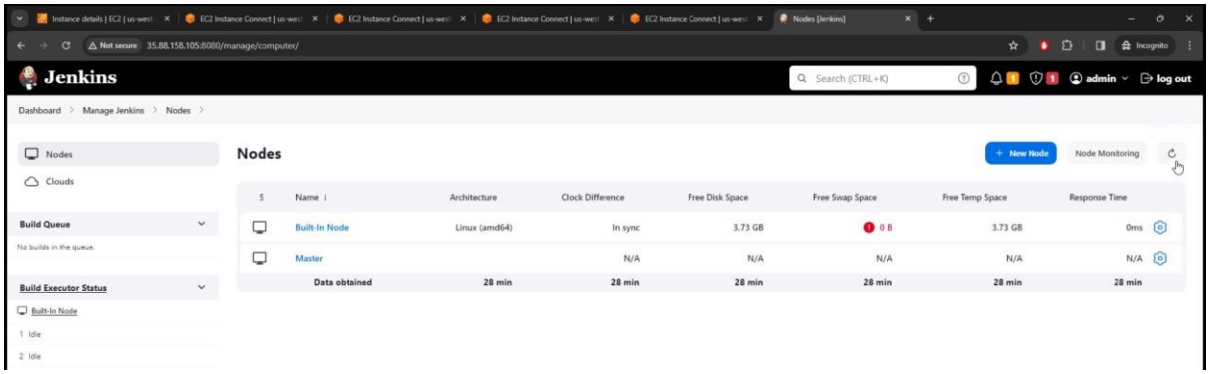

9. Cloning GIT repository and adding Dockerfile, deployment.yaml and service.yaml

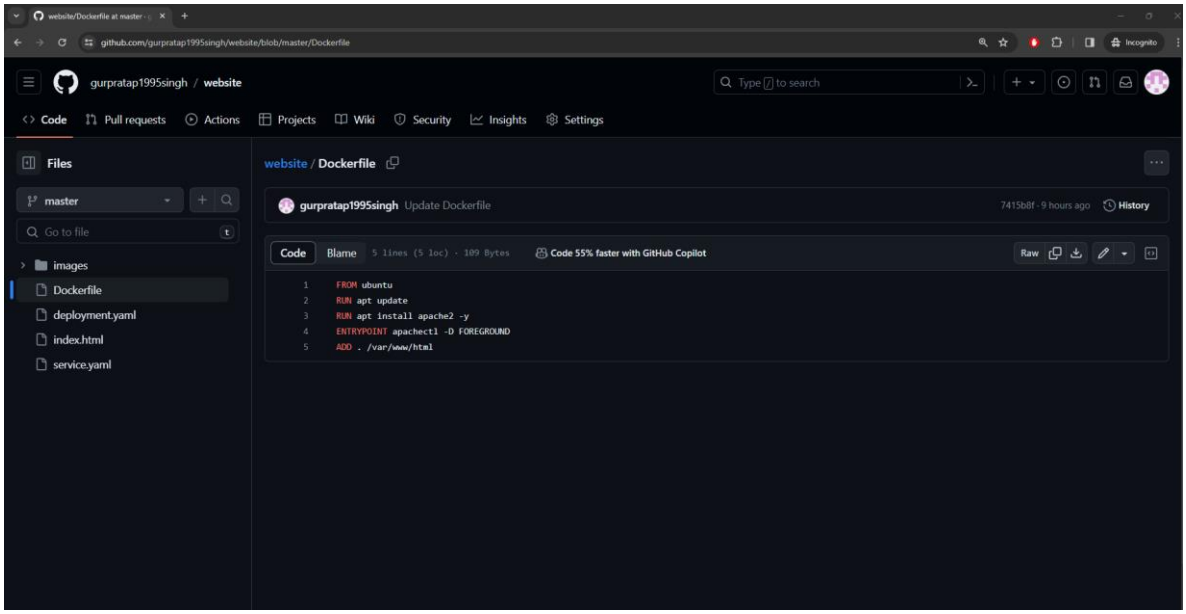

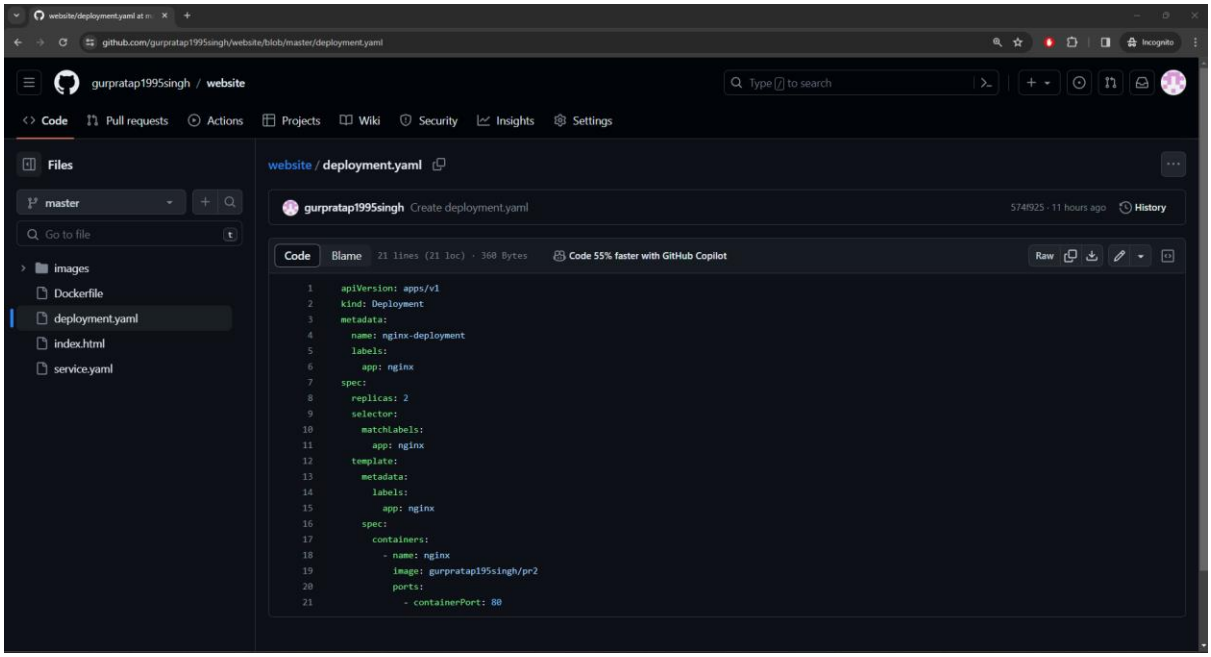

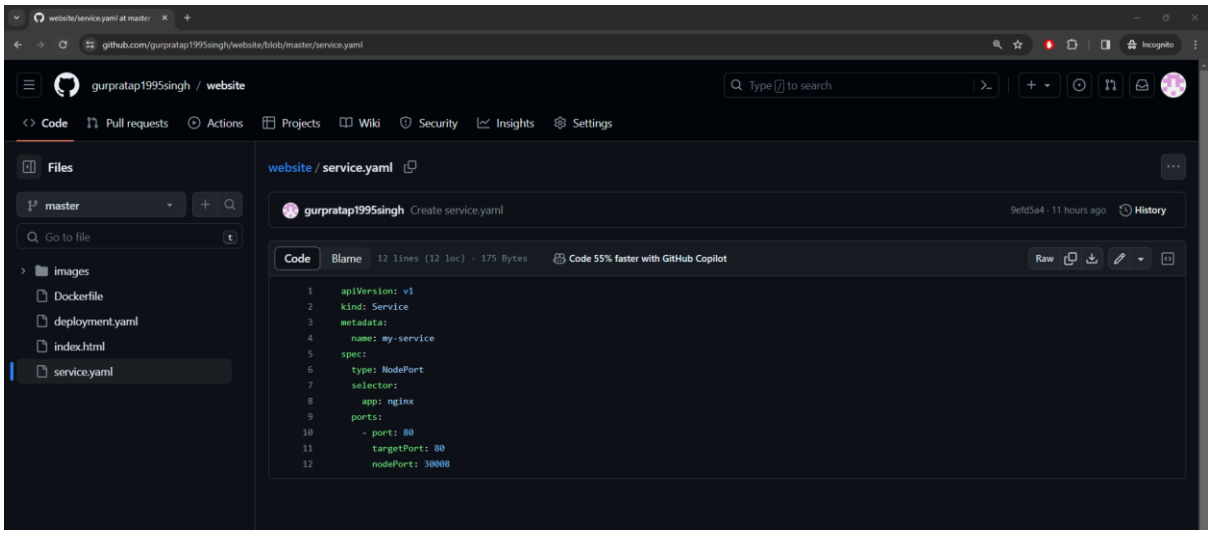

10. Adding Credentials

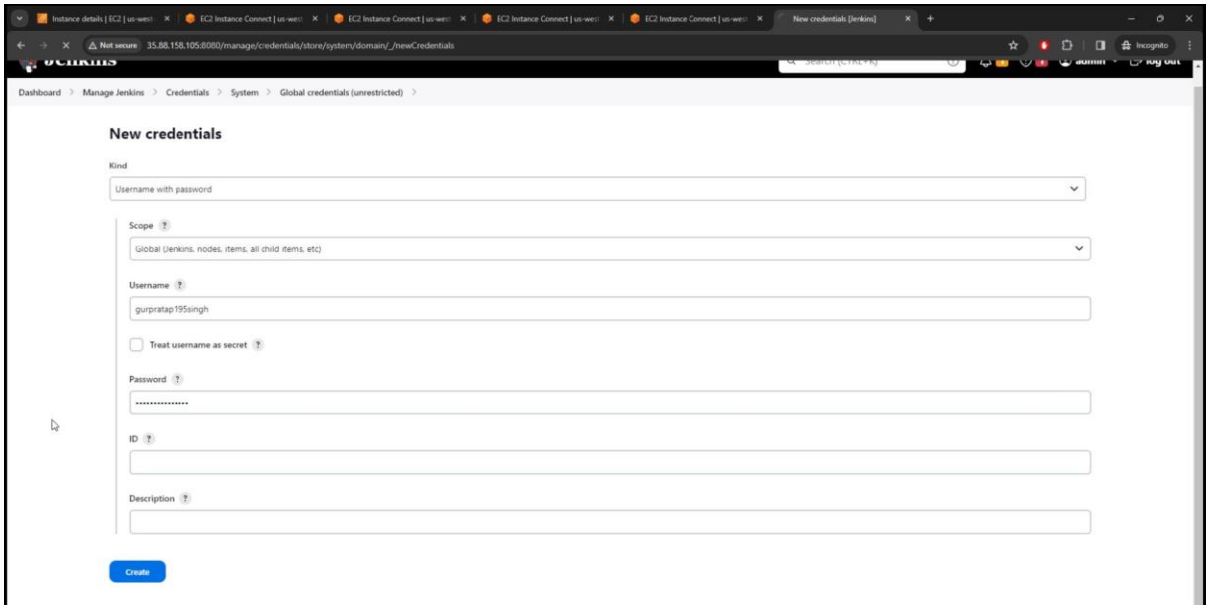

11. Creating Job (Pipeline script).

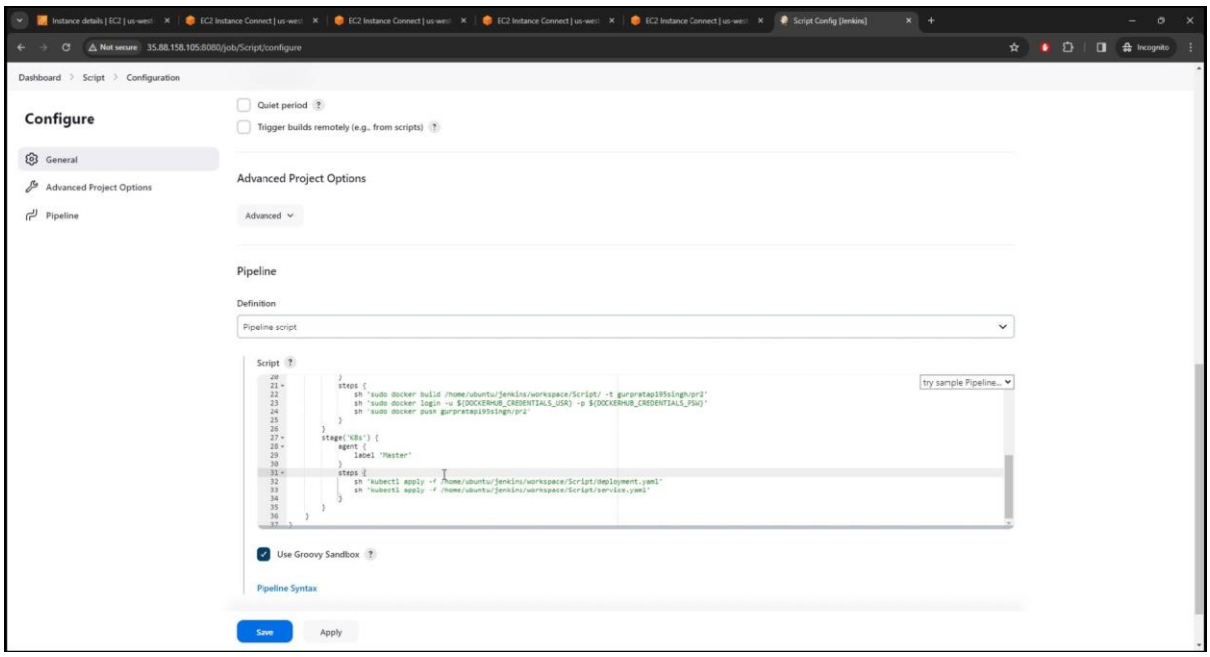

12. Running the Job.

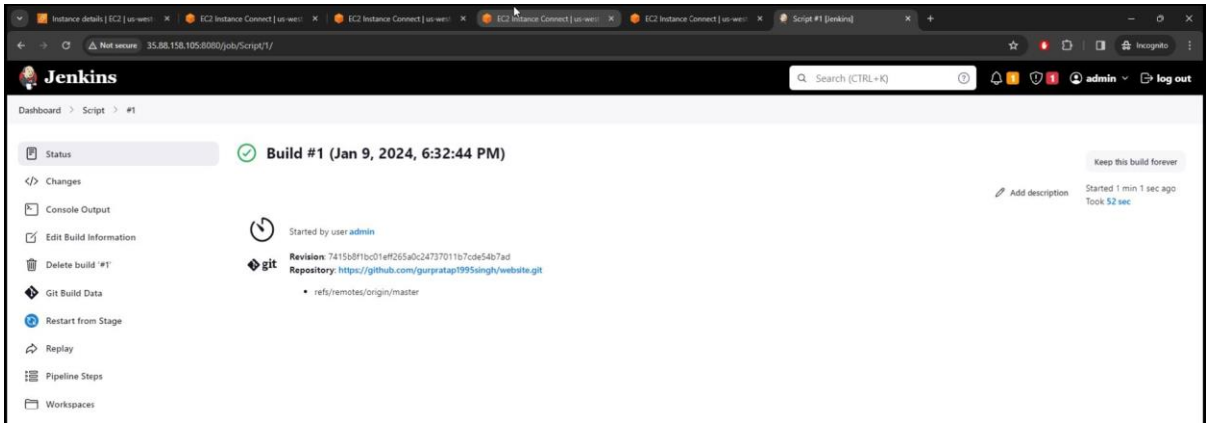

13. Now adding the GitHub Webhook to auto trigger the job.

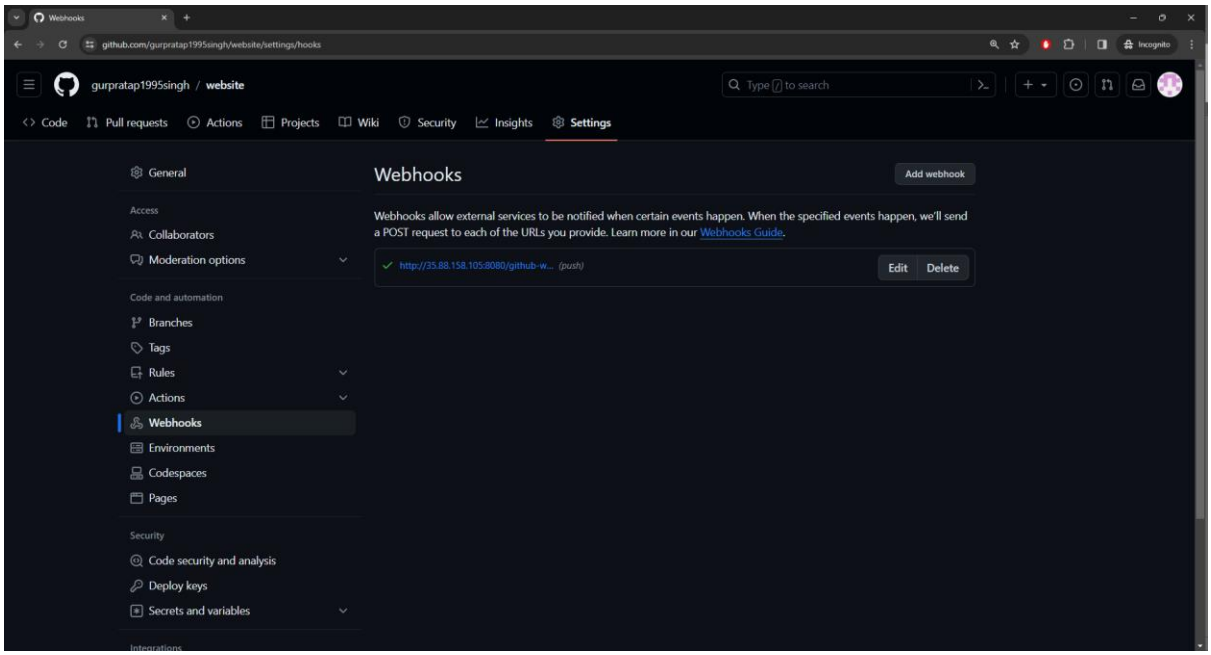

14. Changing in the index.html to verify.

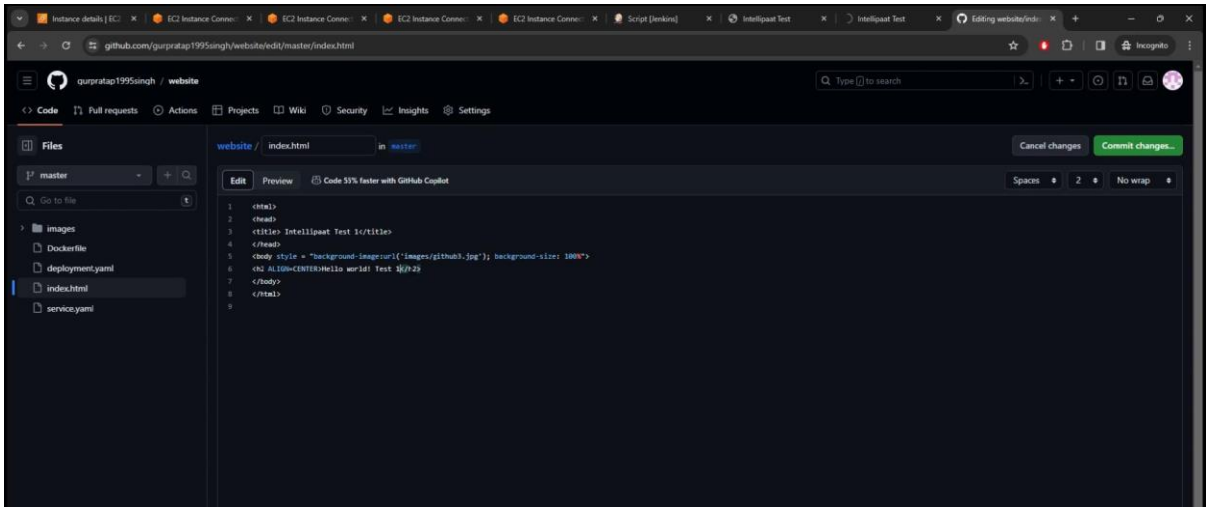

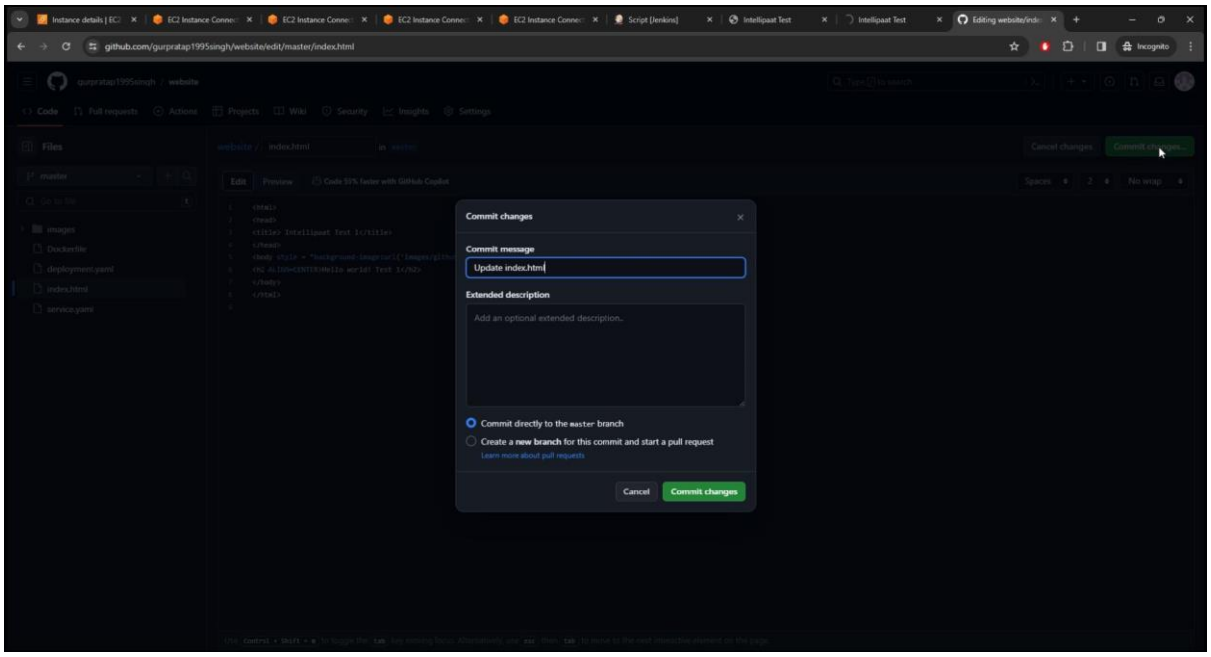
github.com/gurpratap1995singh/website/edit/master/index.html
Commit changes
Commit message
Update index.html
Extended description
Add an optional extended description..
Commit directly to the master branch
Create a new branch for this commit and start a pull request
Cancel
Commit changes

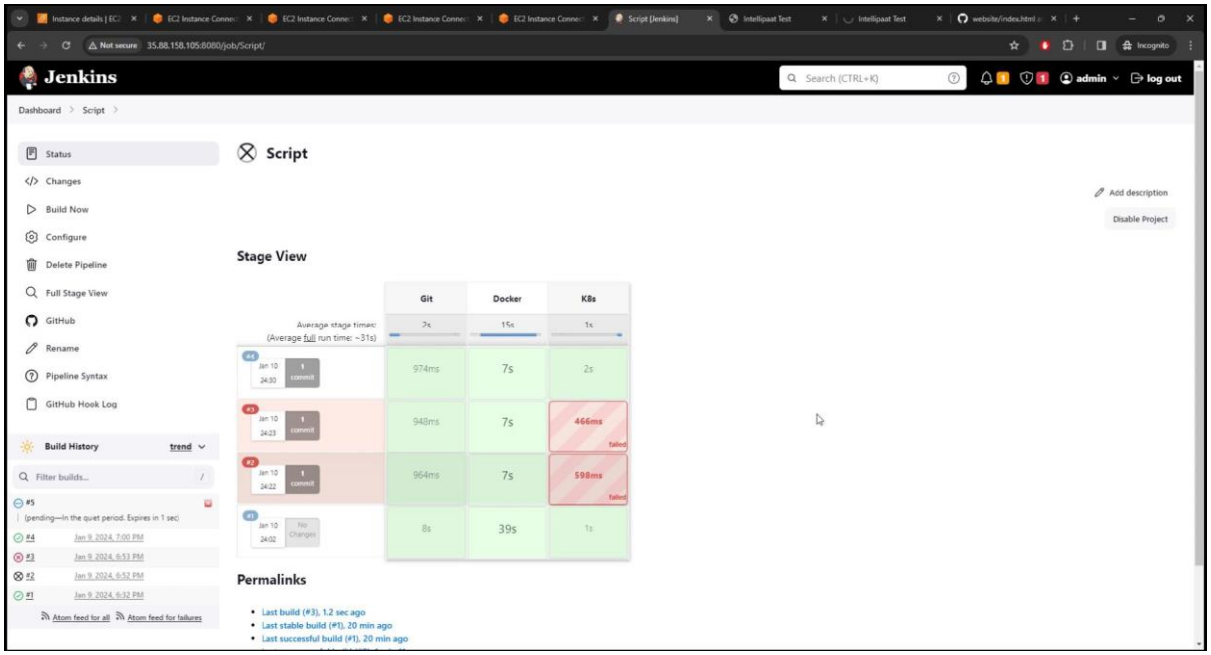
Not secure 35.88.158.105:8080/job/Script/
Jenkins
admin
log out
Dashboard
Script
Status
Changes
Build Now
Configure
Delete Pipeline
Full Stage View
GitHub
Rename
Pipeline Syntax
GitHub Hook Log
Build History
trend
Add description
Disable Project
Stage View
Git
Docker
K8s
Permalinks
Last build (#3), 1.2 sec ago
Last stable build (#1), 20 min ago
Last successful build (#1), 20 min ago

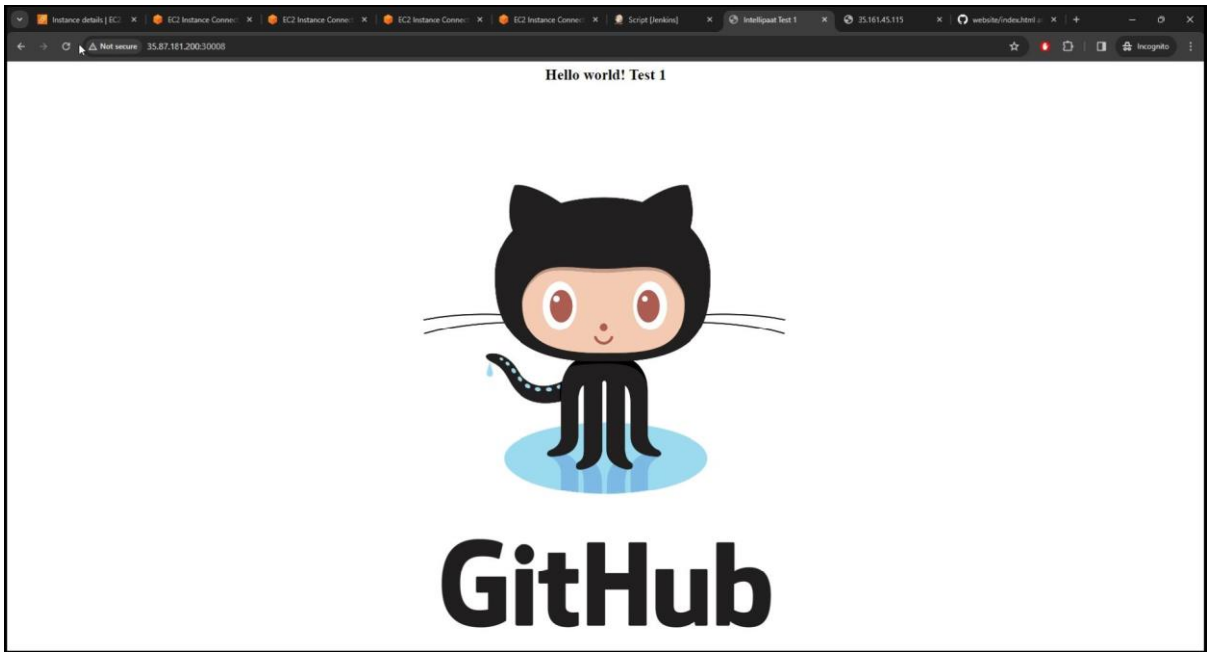
Not secure 35.87.181.200:30008
Incognito
Hello world! Test 1
GitHub

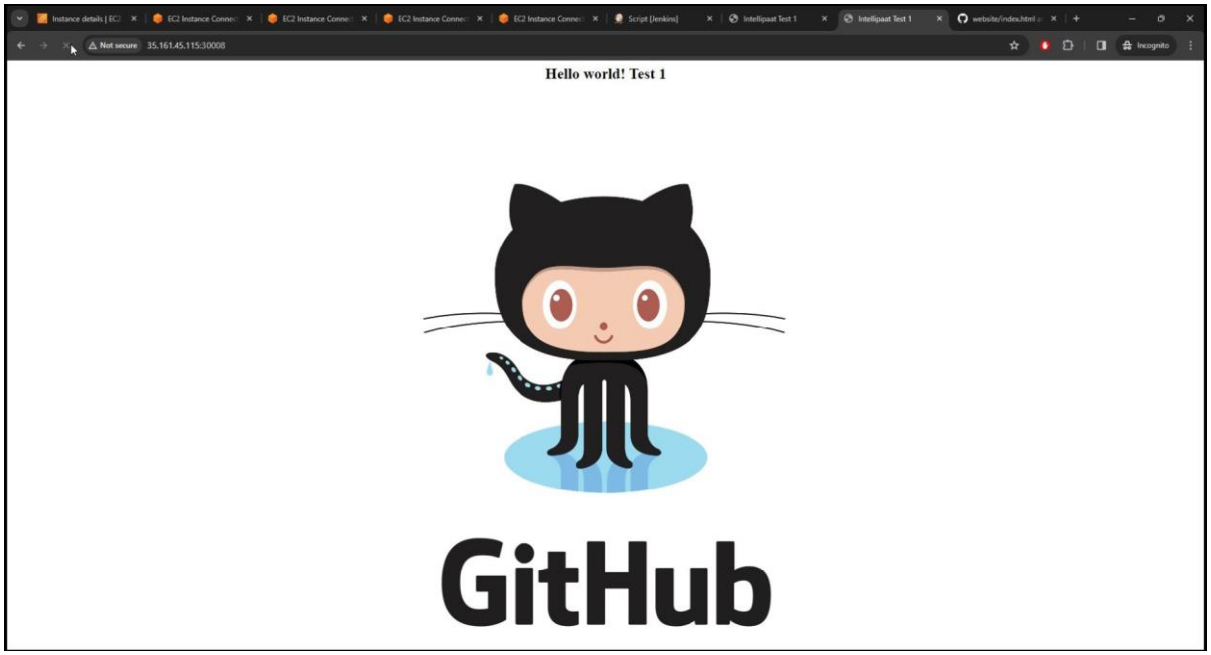
Not secure 35.161.45.115:30008
Incognito
Hello world! Test 1
GitHub